

COMP4900 Assignment 3

YanPeng Gao
101090653

Drew Suitor
101003158

May 21, 2020

Abstract

The task of this project was image classification on images consisting of three separate articles of clothing. The label of the image is associated with the most expensive article of clothing. We were able to show that image preprocessing and fine tuning neural network architectures can improve results and generate an accuracy over 90%.

1 Introduction

The task at hand was to, given an image of 3 separate articles of clothing, classify the image based on the most expensive article of clothing in the image. This classification ranged from 0, the lowest valued item, to 9, the highest valued item. The dataset we are using is a subset of 60,000 data points from the Fashion-MNIST dataset.

There were a handful of approaches we used when assessing this problem. We initially started with an AlexNet, but eventually decided on implementing a customized version of Residual Network as we had better success with it. We will be describing with more depth these approaches later in the report.

The key finding of this report is that we were able to achieve greater than 90 percent accuracy when classifying images from the Fashion-MNIST dataset through the preprocessing of our data and the fine tuning of our implemented Residual Network.

2 Dataset

The dataset used for this image classification task is a modified version of Fashion-MNIST. The dataset consists of 60,000 data points where each point contains an image of 3 articles clothing as well as a corresponding tag ranging from 0 to 9 which corresponds to the value of the most expensive article of clothing in the image.

The images in the dataset contained a lot of empty space between the articles of clothing. We hypothesized that we would see better classification accuracy by preprocessing the images to remove the empty space between the articles of clothing. We took advantage of K-Means clustering to achieve this. K-Means is an algorithm that partitions the dataset, or in our case an individual image, into K non-overlapping subgroups. We clustered on all pixels that were not black pixels to determine the location of individual articles of clothing and then reconstructed the image containing only the 3 articles of clothing.




Figure 1: Image before and after preprocessing




Figure 2: Heatmap of location of icons before preprocessing

We tested out our hypothesis by training and validating our AlexNet model with both the dataset of cropped and uncropped images. We compared the validation accuracy as well as the training loss and found significant improvements on the cropped dataset.




Figure 3: Validation Accuracy and Training Loss on Cropped and Uncropped Datasets

3 Proposed Approach

As with the standard of image classification tasks, convolutional neural networks were used. CNN's have the advantage over regular multilayer perceptrons as they are able to interpret the spatial arrangement of a picture's pixels.

In a CNN, the hidden neuron is made up of input neurons from hidden receptive fields. To do this, a filter (also known as a kernel) of $k_1 \times k_2 \times depth$ is slid across the image. Per each stride, the dot product of the weights in the filter and the neurons in the local receptive field becomes 1 hidden neuron in 1 activation map; and so having multiple features creates multiple activation maps. The purpose of these filters is to activate in the presence of particular shapes, edges, clusters, patterns etc and so the neural network is trying to learn the weights of these filters.

3.1 AlexNet

Alexnet (Krizhevsky et al. [4]) was the first CNN architecture implemented in this project. It was an 8 layer network with 5 convolutional layers and 3 fully connected layers using RELU activation functions instead of sigmoids. While the validation accuracy after 200 epochs was around 93%, the testing accuracy on Kaggle was only 85%. We suspect, as each datapoint was constructed of 3 fashion mnist icons, there were duplication of icons in both the training and validation set hence leading to data leakage. Because of this, the validation set accuracy in these experiments may not generalize well to the testing accuracy.

3.2 Residual Networks

Residual Networks were chosen as they, and variations of ResNets, performed the best in recent years on image classification tasks (Khan et al. [3]). The unique aspect of Residual Networks is the utilization of skip connections every two convolutional layers (He et al. [1]).




Figure 4: Residual Block




Figure 5: Detailed View of Steps in a Residual Block

The residual block in Figure 3 shows the concept of the skip connection. Fundamentally, we are still learning an underlying hypothesis mapping $H(x)$. However, instead of learning $H(x)$ directly, we learn the residual $F(x)$ and have $H(x) = F(x) + x$. The skip connection x is the identity mapping and $F(x)$ is the residual of $H(x) - x$. It is hypothesized in the original paper that it is easier to learn just the residual mapping than the original unreferenced mapping.

The problem with vanishing and exploding gradients makes deeper neural networks harder to train (Pascanu et al. [5]). Without residual blocks, one may encounter that the output mapping of a deeper convolutional layer retains less knowledge than its input. But with the skip connection, we ensure that we always retain the mappings of the previous layer. Therefore each successive layer is at least as good as the last but potentially could learn additional information as well.

The behaviour of residual networks closely resembles an ensemble of shallower neural networks (Veit et al. [6]). Just like removing one model from an ensemble does not worsen the overall accuracy by much, removing single residual layers has no significant impact on accuracy.

Also in the residual blocks, before the RELU activation function, batch normalization is first performed. Batch normalization normalizes each mini-batch of convolutional outputs with a linear transformation: $y_i = \gamma \hat{x}_i + \beta$ where \hat{x}_i is the unit normalized feature and γ and β are learned features. Batch normalization helps solve the problem known as Internal Covariate Shift (Ioffe and Szegedy [2]) where the input distribution to layers deep in the network changes per mini-batch. With the elimination of the covariate shift, the model is able to use larger learning rates for quicker training.

3.3 Random Erasing

In terms of data augmentation, an approach known as Random Erasing (Zhong et al. [7]) was used in all models. With random erasing, images would have the possibility of a random rectangular portion of the image omitted from training. The authors of the paper has shown partial occlusion has the ability to reduce the variance and improve the generalization ability of CNN's.

Our images are very noisy due to the fact that of the three fashion icons per image, possibly only one of them is relevant in the decision making of our CNN. Hence adding random erasing gives the network the ability to encounter instances where some of the noise is blocked out, leading to better learning.

4 Tuning Our Residual Network

It would not make sense to directly use the residual network model presented in PyTorch as it was built for training on $224 \times 224 \times 3$ CIFAR images. The initial architecture style was taken from Zhong et al. [7]'s work as they also produced results for Fashion Mnist (unmodified). Given a depth d and $n = (d - 2)/6$, they proposed an initial convolution layer of 16 filters with a stride of 1 (same convolution), followed by 3 groups of n residual blocks with 16, 32 and 64 filters per convolution respectively. The last 2 groups of residual blocks also halve the length and width by increasing the stride length to 2.

Ideally a grid search of all possible combinations would ensure the optimal network but given our time and computational restraints, the plan is to optimize one parameter, than move on to the next.

4.1 Optimal Depth

In the first step, we compare depths of 20, 32 and 44 layers to find the optimal depth. The criteria were the training loss of 50,000 training images, validation accuracy on 10,000 validation images and the training time per epoch. After this experimentation, a depth of 32 was chosen as it produced the lowest training loss, the highest validation accuracy and was the 2nd quickest to train.




Figure 6: Investigating depth of residual network

4.2 Optimal Filters

We compared groups of blocks using [16,16,32,64] filters per convolutional layer as described above and using more filters of [32,64,128,256] for groups of residual blocks. After experimentation, it was seen that increasing the amount of convolutional filters used helped our model.




Figure 7: Investigating the amount of filters

4.3 Second Fully Connected Layer

The last layer of the original resnet is a fully connected layer where the inputs are a global average pool of each final activation map. In our case, the input would have 256 neurons representing the 256 activation maps.

The inspiration of adding a second fully connected layer comes from our own decision making of classifying the modified fashion mnist. The first step humans make is to recognize the icons while the 2nd step is to classify based on the maximum rank of the 3 icons. Therefore, the two step approached can be incorporated into our residual network with two fully connected layers at the end. We see by the training loss that indeed a second fully connected layer performs better. Validation accuracy was not used due to the data leakage problem.

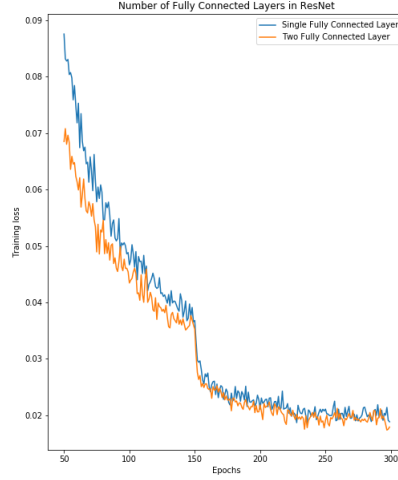


Figure 8: Investigating a second fully connected layer Zhong et al. [7]

5 Results

Initially training was done using 50,000 images whilst a 10,000 set of validation images were used to evaluate accuracy and prevent overfitting. But due to the data leakage problem explained above, a validation set was not deemed able to properly do the two tasks. It was determined to be more useful to train with the full 60,000 training set and use the training loss to evaluate model performance.

Model	Training Loss	Public Test Accuracy
AlexNet UnCropped	0.481	76.466%
AlexNet Cropped	0.239	85.100%
ResNet20- max filters:64	0.117	90.033%
ResNet32- max filters:64	0.108	90.966%
ResNet44- max filters:64	0.125	91.066%
ResNet32- max filters:256, 1FC	0.020	92.36%
ResNet33- max filters:256, 2FC	0.018	92.766%

6 Discussion and Conclusion

The first noticeable jump in our experimentation was by switching from AlexNet to ResNet. We believe that this truly shows the capabilities of skip connections and residual blocks. Image augmentation by using Kmeans to extract the icons helped reduce noise while implementing random erasing improved our models' ability to generalize to the test set. It was learned through experimentation that fine tuning the specific architecture of residual networks to suit your specific classification task also improved results.

Future work can be done on implementing improved versions of ResNets such as DenseNets, PyramidNets or Wide ResNets.

7 Statement of Work

Both YanPeng Gao and Drew Suitor contributed to the report and experimentation. YanPeng Gao also researched and set up the code base for this project.

References

- [1] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. pages 770–778, 06 2016. doi: 10.1109/CVPR.2016.90.
- [2] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. 37:448–456, 07–09 Jul 2015. URL <http://proceedings.mlr.press/v37/ioffe15.html>.
- [3] Riaz Khan, Xiaosong Zhang, Rajesh Kumar, and Emelia Opoku Aboagye. Evaluating the performance of resnet model based on image recognition. 11 2018. doi: 10.1145/3194452.3194461.
- [4] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. pages 1097–1105, 2012. URL <http://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>.
- [5] Razvan Pascanu, Tomas Mikolov, and Y. Bengio. On the difficulty of training recurrent neural networks. *30th International Conference on Machine Learning, ICML 2013*, 11 2012.
- [6] Andreas Veit, Michael Wilber, and Serge Belongie. Residual networks behave like ensembles of relatively shallow networks. *Advances in Neural Information Processing Systems*, 05 2016.
- [7] Zhun Zhong, Liang Zheng, Guoliang Kang, Shaozi Li, and Yi Yang. Random erasing data augmentation. *CoRR*, abs/1708.04896, 2017. URL <http://arxiv.org/abs/1708.04896>.

8 Appendix

8.1 Kmeans Segmentation

```
1 import pickle
2 import matplotlib.pyplot as plt
3 import numpy as np
4 from torchvision import transforms
5 from torch.utils.data import Dataset
6 from torch.utils.data import DataLoader
7 from PIL import Image
8 import torch
9
10 data = pickle.load( open( './data/Test.pkl', 'rb' ), encoding='bytes')
11 targets = np.genfromtxt('./data/TrainLabels.csv', delimiter=',')
12
13
14 from sklearn import cluster
15
16 kmeans_cluster = cluster.KMeans(n_clusters=3, init='k-means++')
17 newdatalist = []
18 for index in range(data.shape[0]):
19     arr1 = np.argwhere(data[index,:,:] > 20) #2d array where each row has x and y
        coordinates of non black values
20     kmeans_cluster.fit(arr1)
21     cluster_centers = kmeans_cluster.cluster_centers_
22     newimagelist = []
23     for i in range(len(cluster_centers)):
24         row = int(round(cluster_centers[i][0]))
25         col = int(round(cluster_centers[i][1]))
26         if row<=13:
27             row = 14
28         if row >= 50:
29             row = 49
30         if col<=13:
31             col = 14
32         if col >= 114:
33             col = 113
34         newimagelist.append(data[index,row-14:row+14,col-14:col+14])
35     newimage = np.concatenate(newimagelist,axis=1)
36     newdatalist.append(newimage)
37 newdata = np.stack(newdatalist)
38 newdata.dump("./data/NewTest.pkl")
```

Listing 1: Image Cropping KMeans

8.2 ResNet

```
1 # -*- coding: utf-8 -*-
2 """ResNet.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1j8RtFmvL0Abmli9sr30i1xZVXamfTewN
8
9 # Introduction
10 In the following you will see how to read the provided files for the mini-project 3.
11 First you will see how to read each of the provided files. Then, you will see a more
12     elegant way of using this data for training neural networks.
13
14 from google.colab import drive
15 drive.mount('/content/gdrive' )
16
17 # Commented out IPython magic to ensure Python compatibility.
18 # %cd '/content/gdrive/My Drive/Comp4900A3'
19 !ls './data/'
20
21 import pickle
22 import matplotlib.pyplot as plt
23 import numpy as np
24 from torchvision import transforms
25 from torch.utils.data import Dataset
26 from torch.utils.data import DataLoader
27 from PIL import Image
28 import torch
29 import math
30
31 # Read a pickle file and display its samples
32 # Note that image data are stored as unit8 so each element is an integer value between
33     0 and 255
34 data = pickle.load( open( './data/NewTrain.pkl', 'rb' ), encoding='bytes')
35 targets = np.genfromtxt('./data/TrainLabels.csv', delimiter=',')
36 plt.imshow(data[1234,:,:], cmap='gray', vmin=0, vmax=256)
37
38 """# Dataset class
39 *Dataset* class and the *Dataloader* class in pytorch help us to feed our own training
40     data into the network. Dataset class is used to provide an interface for accessing
41     all the training or testing samples in your dataset. For your convinance, we
42     provide you with a custom Dataset that reads the provided data including images (.
43     pkl file) and labels (.csv file).
44
45 # Dataloader class
46 Although we can access all the training data using the Dataset class, for neural
47     networks, we would need batching, shuffling, multiprocessing data loading, etc.
48     DataLoader class helps us to do this. The DataLoader class accepts a dataset and
49     other parameters such as batch_size.
50
51 """
52
53 # Transforms are common image transformations. They can be chained together using
54     Compose.
55 # Here we normalize images img=(img-0.5)/0.5
56 img_transform = transforms.Compose([
57     transforms.RandomHorizontalFlip(),
58     transforms.ToTensor(),
59     transforms.Normalize([0.5], [0.5]),
60     transforms.RandomErasing(p=0.5, ratio=(0.5,2.0)),
61 ])
62
63 test_transform = transforms.Compose([
```



```

54     transforms.ToTensor(),
55     transforms.Normalize([0.5], [0.5])
56 ])
57
58 class MyDataset(Dataset):
59     def __init__(self, img_file, label_file, transform=None, idx = None):
60         self.data = pickle.load( open( img_file, 'rb' ), encoding='bytes')
61         self.targets = np.genfromtxt(label_file, delimiter=',')
62         if idx is not None:
63             self.targets = self.targets[idx]
64             self.data = self.data[idx]
65             self.transform = transform
66
67     def __len__(self):
68         return len(self.targets)
69
70     def __getitem__(self, index):
71         img, target = self.data[index], int(self.targets[index])
72         img = Image.fromarray(img.astype('uint8'), mode='L')
73
74         if self.transform is not None:
75             img = self.transform(img)
76
77         return img, target
78
79 # Read image data and their label into a Dataset class
80 dataset = MyDataset('./data/NewTrain.pkl', './data/TrainLabels.csv', transform=
    img_transform, idx=None)
81 train_set, val_set = torch.utils.data.random_split(dataset, [59999, 1])
82
83 batch_size = 256 #feel free to change it
84 trainloader = DataLoader(train_set, batch_size=batch_size, shuffle=True)
85
86 device = 'cuda' #CUDA is GPU
87
88 import torch.nn as nn
89 import torch.nn.functional as F
90 #code from:
91 # https://github.com/pytorch/vision/blob/master/torchvision/models/resnet.py
92 # https://github.com/zhunzhong07/Random-Erasing/blob/master/models/fashion/resnet.py
93
94 def conv3x3(in_planes, out_planes, stride=1):
95     "3x3 convolution with padding"
96     return nn.Conv2d(in_planes, out_planes, kernel_size=3, stride=stride,
97                     padding=1, bias=False)
98
99 class Bottleneck(nn.Module):
100     expansion = 4
101
102     def __init__(self, inplanes, planes, stride=1, downsample=None):
103         super(Bottleneck, self).__init__()
104         self.conv1 = nn.Conv2d(inplanes, planes, kernel_size=1, bias=False)
105         self.bn1 = nn.BatchNorm2d(planes)
106         self.conv2 = nn.Conv2d(planes, planes, kernel_size=3, stride=stride,
107                             padding=1, bias=False)
108         self.bn2 = nn.BatchNorm2d(planes)
109         self.conv3 = nn.Conv2d(planes, planes * 4, kernel_size=1, bias=False)
110         self.bn3 = nn.BatchNorm2d(planes * 4)
111         self.relu = nn.ReLU(inplace=True)
112         self.downsample = downsample
113         self.stride = stride
114
115     def forward(self, x):
116         residual = x
117

```

```

118         out = self.conv1(x)
119         out = self.bn1(out)
120         out = self.relu(out)
121
122         out = self.conv2(out)
123         out = self.bn2(out)
124         out = self.relu(out)
125
126         out = self.conv3(out)
127         out = self.bn3(out)
128
129         if self.downsample is not None:
130             residual = self.downsample(x)
131
132         out += residual
133         out = self.relu(out)
134
135         return out
136
137 class BasicBlock(nn.Module):
138     expansion = 1
139
140     def __init__(self, inplanes, planes, stride=1, downsample=None):
141         super(BasicBlock, self).__init__()
142         self.conv1 = conv3x3(inplanes, planes, stride)
143         self.bn1 = nn.BatchNorm2d(planes)
144         self.relu = nn.ReLU(inplace=True)
145         self.conv2 = conv3x3(planes, planes)
146         self.bn2 = nn.BatchNorm2d(planes)
147         self.downsample = downsample
148         self.stride = stride
149
150     def forward(self, x):
151         residual = x
152
153         out = self.conv1(x)
154         out = self.bn1(out)
155         out = self.relu(out)
156
157         out = self.conv2(out)
158         out = self.bn2(out)
159
160         if self.downsample is not None:
161             residual = self.downsample(x)
162
163         out += residual
164         out = self.relu(out)
165
166         return out
167
168 class ResNet(nn.Module):
169
170     def __init__(self, depth=20, num_classes=10):
171         super(ResNet, self).__init__()
172         # Model type specifies number of layers for CIFAR-10 model
173         assert (depth - 2) % 6 == 0, 'depth should be 6n+2'
174         n = int((depth - 2) / 6)
175
176         block = Bottleneck if depth >= 44 else BasicBlock
177
178         self.inplanes = 32
179         self.conv1 = nn.Conv2d(1, self.inplanes, kernel_size=3, padding=1,
180                                bias=False)
181         self.bn1 = nn.BatchNorm2d(self.inplanes)
182         self.relu = nn.ReLU(inplace=True)

```

```

183 self.layer1 = self._make_layer(block, 64, n) #maybe this should have been 32
184 self.layer2 = self._make_layer(block, 128, n, stride=2)
185 self.layer3 = self._make_layer(block, 256, n, stride=2)
186 # self.avgpool = nn.AvgPool2d(7)
187 self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
188 self.fc1 = nn.Linear(256 * block.expansion, 220)
189 self.fc2 = nn.Linear(220, num_classes)
190
191 for m in self.modules():
192     if isinstance(m, nn.Conv2d):
193         n = m.kernel_size[0] * m.kernel_size[1] * m.out_channels
194         m.weight.data.normal_(0, math.sqrt(2. / n))
195     elif isinstance(m, nn.BatchNorm2d):
196         m.weight.data.fill_(1)
197         m.bias.data.zero_()
198
199 def _make_layer(self, block, planes, blocks, stride=1):
200     downsample = None
201     if stride != 1 or self.inplanes != planes * block.expansion:
202         downsample = nn.Sequential(
203             nn.Conv2d(self.inplanes, planes * block.expansion,
204                 kernel_size=1, stride=stride, bias=False),
205             nn.BatchNorm2d(planes * block.expansion),
206         )
207
208     layers = []
209     layers.append(block(self.inplanes, planes, stride, downsample))
210     self.inplanes = planes * block.expansion
211     for i in range(1, blocks):
212         layers.append(block(self.inplanes, planes))
213
214     return nn.Sequential(*layers)
215
216 def forward(self, x):
217     x = self.conv1(x)
218     x = self.bn1(x)
219     x = self.relu(x)
220
221     x = self.layer1(x)
222     x = self.layer2(x)
223     x = self.layer3(x)
224
225     x = self.avgpool(x)
226     x = x.view(x.size(0), -1)
227     x = self.fc1(x)
228     x = self.fc2(x)
229     return x
230
231 net = ResNet(num_classes=10, depth=32)
232 net.to(device)
233 import torch.optim as optim
234
235 criterion = nn.CrossEntropyLoss()
236
237 optimizer = optim.SGD(net.parameters(), lr=0.1, momentum=0.9)
238
239 # valloader = DataLoader(val_set, batch_size=batch_size, shuffle=True)
240 trainingloss = []
241 for epoch in range(400): # loop over the dataset multiple times
242     net.train()
243     running_loss = 0.0
244
245     for i, data in enumerate(trainloader, 0):
246
247         # get the inputs; data is a list of [inputs, labels]

```

```

248     inputs, labels = data[0].to(device), data[1].to(device)
249
250     # zero the parameter gradients
251     optimizer.zero_grad()
252
253     # forward + backward + optimize
254     outputs = net(inputs)
255     loss = criterion(outputs, labels)
256     loss.backward()
257     optimizer.step()
258     # print statistics
259     running_loss += loss.item()
260
261     print('[%d] training loss: %.5f' % (epoch + 1, running_loss / i))
262     trainingloss.append(running_loss / i)
263
264     if epoch == 99:
265         PATH = './saved_models/resnet6_100epochs.pth'
266         torch.save(net.state_dict(), PATH)
267     if epoch == 199:
268         PATH = './saved_models/resnet6_200epochs.pth'
269         torch.save(net.state_dict(), PATH)
270     if epoch == 299:
271         PATH = './saved_models/resnet6_300epochs.pth'
272         torch.save(net.state_dict(), PATH)
273     if epoch == 399:
274         PATH = './saved_models/resnet6_400epochs.pth'
275         torch.save(net.state_dict(), PATH)
276
277
278     if epoch == 150:
279         for param_group in optimizer.param_groups:
280             param_group['lr'] = 0.01
281     if epoch == 225:
282         for param_group in optimizer.param_groups:
283             param_group['lr'] = 0.001
284     if epoch == 300:
285         for param_group in optimizer.param_groups:
286             param_group['lr'] = 0.0005
287
288     print('Finished Training')
289
290
291
292 testdata = MyDataset('./data/NewTest.pkl', './data/temptestlabels.csv', transform=
    test_transform, idx=None)
293 testloader = DataLoader(testdata, batch_size=batch_size, shuffle=False)
294
295 tempdata = pickle.load( open( './data/NewTest.pkl', 'rb' ), encoding='bytes')
296 print(tempdata.shape)
297 plt.imshow(tempdata[5555], cmap='gray', vmin=0, vmax=256)
298
299 correct = 0
300 total = 0
301 loopcount=0
302 preds = []
303 net.eval()
304 with torch.no_grad():
305     for data in testloader:
306         images, labels = data[0].to(device), data[1].to(device)
307         outputs = net(images)
308         _, predicted = torch.max(outputs, 1)
309         preds.append(predicted.cpu().numpy())
310         total += labels.size(0)
311         correct += (predicted == labels).sum().item()

```

```

312         loopcount+=1
313     print('Accuracy of the network on the 10000 test images: %d %%' % (
314         100 * correct / total))
315
316     import numpy as np
317     preds = np.concatenate(preds,axis=0)
318
319     import pandas as pd
320     indexlist = list(range(0,10000))
321     df = pd.DataFrame(list(zip(indexlist, preds)),
322                       columns =['id', 'output'])
323
324     df.to_csv("./Submissions/resnet34morefilters_200epochs.csv",index=False)
325
326     #ResNet

```

Listing 2: ResNet

8.3 AlexNet

```
1 #AlexNet
2 img_transform = transforms.Compose([
3     transforms.ToTensor(),
4     transforms.Normalize([0.5], [0.5])
5 ])
6 class MyDataset(Dataset):
7     def __init__(self, img_file, label_file, transform=None, idx = None):
8         self.data = pickle.load( open( img_file, 'rb' ), encoding='bytes')
9         self.targets = np.genfromtxt(label_file, delimiter=',')
10        if idx is not None:
11            self.targets = self.targets[idx]
12            self.data = self.data[idx]
13        self.transform = transform
14
15    def __len__(self):
16        return len(self.targets)
17
18    def __getitem__(self, index):
19        img, target = self.data[index], int(self.targets[index])
20        img = Image.fromarray(img.astype('uint8'), mode='L')
21
22        if self.transform is not None:
23            img = self.transform(img)
24
25        return img, target
26
27 dataset = MyDataset('./data/NewTrain.pkl', './data/TrainLabels.csv', transform=
    img_transform, idx=None)
28 old_dataset = MyDataset('./data/Train.pkl', './data/TrainLabels.csv', transform=
    img_transform, idx=None)
29
30 train_set, val_set = torch.utils.data.random_split(dataset, [50000, 10000])
31 old_train_set, old_val_set = torch.utils.data.random_split(old_dataset, [50000,
    10000])
32
33 batch_size = 256 #feel free to change it
34 trainloader = DataLoader(train_set, batch_size=batch_size, shuffle=True)
35 old_trainloader = DataLoader(old_train_set, batch_size=batch_size, shuffle=True)
36
37 import torch.nn as nn
38 import torch.nn.functional as F
39
40
41 class AlexNet(nn.Module):
42
43     def __init__(self, num_classes=1000):
44         super(AlexNet, self).__init__()
45         self.features = nn.Sequential(
46             nn.Conv2d(1, 64, kernel_size=6, stride=2, padding=2),
47             nn.ReLU(inplace=True),
48             nn.MaxPool2d(kernel_size=2, stride=2),
49             nn.Conv2d(64, 192, kernel_size=5, padding=2),
50             nn.ReLU(inplace=True),
51             nn.MaxPool2d(kernel_size=3, stride=2),
52             nn.Conv2d(192, 384, kernel_size=3, padding=1),
53             nn.ReLU(inplace=True),
54             nn.Conv2d(384, 256, kernel_size=3, padding=1),
55             nn.ReLU(inplace=True),
56             nn.Conv2d(256, 256, kernel_size=3, padding=1),
57             nn.ReLU(inplace=True),
58             nn.MaxPool2d(kernel_size=3, stride=2),
59         )
60         self.avgpool = nn.AdaptiveAvgPool2d((6, 6))
```

```

61         self.classifier = nn.Sequential(
62             nn.Dropout(),
63             nn.Linear(256 * 6 * 6, 4096),
64             nn.ReLU(inplace=True),
65             nn.Dropout(),
66             nn.Linear(4096, 4096),
67             nn.ReLU(inplace=True),
68             nn.Linear(4096, num_classes),
69         )
70
71     def forward(self, x):
72         x = self.features(x)
73         x = self.avgpool(x)
74         x = torch.flatten(x, 1)
75         x = self.classifier(x)
76         return x
77
78
79 net = AlexNet(num_classes=10)
80 net.to(device)
81 import torch.optim as optim
82
83 criterion = nn.CrossEntropyLoss()
84 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
85
86
87 valloader = DataLoader(val_set, batch_size=batch_size, shuffle=True)
88 training_loss = [0] * 100
89 validation_acc = [0] * 100
90 for epoch in range(100): # loop over the dataset multiple times
91
92     running_loss = 0.0
93     for i, data in enumerate(trainloader, 0):
94         # get the inputs; data is a list of [inputs, labels]
95         inputs, labels = data[0].to(device), data[1].to(device)
96
97         # zero the parameter gradients
98         optimizer.zero_grad()
99
100        # forward + backward + optimize
101        outputs = net(inputs)
102        loss = criterion(outputs, labels)
103        loss.backward()
104        optimizer.step()
105        # print statistics
106        running_loss += loss.item()
107
108    training_loss[epoch] = (running_loss/i)
109    print('[%d] training loss: %.5f' % (epoch + 1, running_loss / i))
110
111    correct = 0
112    total = 0
113    with torch.no_grad():
114        for data in valloader:
115            images, labels = data[0].to(device), data[1].to(device)
116            outputs = net(images)
117            _, predicted = torch.max(outputs, 1)
118            # _, predicted = torch.topk(input=outputs,k=3,dim=1)
119            # predicted,_ = torch.max(predicted,dim=1)
120            total += labels.size(0)
121            correct += (predicted == labels).sum().item()
122
123    validation_acc[epoch] = (100 * correct/total)
124    print('[%d] Validation acc: %.5f' % (epoch + 1, (100 * correct / total)))
125

```

```
126 print(training_loss)
127 print(validation_acc)
128
129 print('Finished Training')
```

Listing 3: AlexNet